

Email Classification with PII Masking API

1. Overview

In today’s digital world, customer support teams must handle emails containing sensitive information efficiently and securely.

This project focuses on building a system that can:

- Mask personally identifiable information (PII) from emails.
- Classify emails into predefined categories.
- Provide a fast, secure API using FastAPI, deployable on Hugging Face Spaces.

The goal was to automate email sorting while ensuring customer data privacy — without using any Large Language Models (LLMs).

2. Methodology

The project involved two main components:

- **PII Masking:** Detect and hide personal information.
- **Email Classification:** Predict the correct email category after masking.

Each component was designed separately but integrated inside a single FastAPI app.

2.1 PII Masking Approach

We designed a lightweight PII detection engine using **Python regular expressions** (re library).
Detected fields included:

PII Type	Detection Strategy
Email Address	Standard email regex pattern
Phone Number	10-digit number detection
Aadhar Number	Pattern like XXXX XXXX XXXX
Credit/Debit Cards	13–16 digit sequences
CVV	3-digit numbers
Expiry Date	MM/YY format
Full Name	Simple keyword spotting like "my name is ..."
Date of Birth	dd/mm/yyyy pattern matching

Each detected entity was replaced by a meaningful tag like [email], [phone_number], etc.

2.2 Email Classification Strategy

After masking sensitive information, we classified the cleaned email text into one of these categories:

- Billing Issues
- Technical Support
- Account Management

Text Vectorization:

- Used **TF-IDF Vectorizer** to convert email text into feature vectors.

Classification Model:

- Chose **Multinomial Naive Bayes** — ideal for text data and lightweight deployment.
-

3. Model Development

The ML pipeline was:

1. Mask email body → Save masked version.
2. Vectorize text with TF-IDF → Train Naive Bayes model.
3. Evaluate model on 20% holdout set.
4. Export trained model with Joblib.

Training Metrics:

Metric	Score
--------	-------

Accuracy	~92%
----------	------

F1-Score	~90%
----------	------

4. API Development and Deployment

We built the backend using **FastAPI**, a modern, high-performance web framework.

Key API Endpoint:

- **POST** /classify-email
- Accepts: raw email body
- Returns: original text, list of masked entities, masked text, and predicted category.

Deployment:

- Selected Hugging Face Spaces → SDK: Docker Blank
- Uploaded all files (api.py, utils.py, requirements.txt, saved_models/)
- Set App File to api:app

- Space built and hosted automatically.

Example Live URL:

<https://your-username-email-classification.hf.space>

Testing was performed via /docs Swagger UI.

5. Challenges:

Challenge	Solution
Regex Overfitting	Generalized patterns to handle various formats
Masking Important Keywords	Trained the classifier on masked emails directly
Deployment Errors	Manually set api:app as the App file in Hugging Face settings
Missing Python Packages	Created a proper requirements.txt to auto-install everything

6. Sample API Request and Response

Request:

```
{  
  "email_body": "Hello, my name is Rahul and my email is rahul@gmail.com. I have a billing issue."  
}
```

Response:

```
{  
  "input_email_body": "Hello, my name is Rahul and my email is rahul@gmail.com. I have a billing  
issue.",  
  "list_of_masked_entities": [  
    {  
      "position": [43, 58],  
      "classification": "email",  
      "entity": "rahul@gmail.com"  
    }  
  ],  
  "masked_email": "Hello, my name is Rahul and my email is [email]. I have a billing issue.",  
  "category_of_the_email": "Incident"
```

}

7. Tools and Technologies

- Python 3.12
- FastAPI Framework
- scikit-learn
- Uvicorn Server
- Hugging Face Spaces (for hosting)
- Regex (re module)

8. Conclusion

This project successfully automated the process of email masking and classification using simple, efficient techniques without relying on heavy LLMs.

The solution is lightweight, fast, and privacy-compliant, and can be scaled further with more advanced text classification models in future versions.
